Antaeus GTM OS: Architecture Restructure Research Brief

The winning architecture for a founder-led sales operating system is object-centric, behaviorally-engineered, and command-first — and research across 13 competitive tools and 14 psychological principles confirms that Antaeus's four-layer model (Command → Sheet → Workspace → Graph) maps precisely to how expert users naturally process information. The critical insight: Linear, Superhuman, and Attio all converge on the same pattern — a ranked action surface that dissolves module boundaries and treats every entity as an inspectable, actionable object. The behavioral science is equally decisive: the Zeigarnik Effect's motivational drive (not its memory component, which a 2025 meta-analysis debunked) Wikipedia nature combined with the Endowed Progress Effect and Implementation Intentions creates a loop where tension migrates forward through object states rather than resolving at completion. This research synthesizes over 40 sources into actionable architecture decisions for Antaeus's final pre-beta restructure.

The modern B2B UX stack converges on five patterns

Across Linear, Superhuman, Attio, Clay, Notion, Arc, and Raycast, a clear architectural consensus has emerged — one that directly validates and refines Antaeus's planned restructure.

Pattern 1: Command palette as navigation layer, not feature. Raycast proved a command palette can replace an entire OS interaction model. Maggie Appleton Ai Review Superhuman's ⌘K trains its own obsolescence by showing keyboard shortcuts alongside results Superhuman Product Hunt — users graduate from palette to direct shortcuts, creating compounding muscle memory. Blake Crosley superhuman Linear's palette is context-aware, surfacing cycle commands when viewing cycles and issue commands when viewing issues. Linear The consensus: **⌘K should be the universal entry point for Brief/Grid/Queue mode switching, object search, and action execution.** Retool's taxonomy matters here — command palettes are either search-first (Notion, Spotlight) or action-first (Superhuman, Linear). Retool Antaeus should be action-first: "Log call with Acme," "Move deal to Negotiation," "Show today's brief."

Pattern 2: Views are lenses, not modules. Attio's most radical insight is that a pipeline is just a list displayed as kanban, CRM.org and a report is just a list with grouping. Same data, different projections. Workflow Automation This eliminates the traditional CRM navigation problem of "where does this data live?" by making the answer always "in the object." Linear echoes this with display options that let any issue list switch between list, board, and timeline views while maintaining identical underlying data. Linear **For Antaeus, this means the 16 current modules should not map 1:1 to 16 navigation items — they should dissolve into view modes on shared object types.**

Pattern 3: Queue/triage as the primary action model. Superhuman's core innovation converts email from a random-access list into a serial queue. Each item requires exactly one decision: Reply, Archive, Snooze, or Forward. (Arman) Linear's Triage inbox does the same for issues, with AI-suggested assignments, labels, and duplicate detection. (Linear) The psychology is clear: **pre-categorized queues dramatically reduce decision fatigue compared to unified lists** (Hick's Law predicts this — decision time increases logarithmically with options). (LogRocket +2) Antaeus's Queue mode should force a disposition decision on every item: Act, Snooze, Delegate, or Dismiss.

Pattern 4: AI embedded in workflow, not bolted on. Superhuman's Auto Summarize is "the first AI feature you don't have to remember to use" — a one-line summary above every conversation, always on, requiring no invocation. (Superhuman) Attio's Ask Attio operates on the entire CRM graph, accessible from sidebar, home, and record pages. (Attio) (Attio) Clay embeds AI at every column of its enrichment tables. (Skaled +2) Salesforce and HubSpot treat AI as premium add-ons, gated behind additional licensing. (Supportyourapp) **The winning pattern is ambient AI that pre-processes, pre-summarizes, and pre-suggests without explicit invocation.** For Antaeus, AI should generate the Brief narrative, suggest next actions on Sheets, and surface cross-object patterns — all without the user asking.

Pattern 5: Speed as product feature. Linear targets sub-100ms interactions (Onehorizon) and benchmarks **3.7x faster than Jira**. (Onehorizon) Superhuman's internal target is 50-60ms, with key repeat rates of 65ms versus macOS's default 100ms. (blakecrosley) Both use optimistic UI — actions complete visually before server confirmation, with undo as the safety net. (blakecrosley) **Perceived speed IS the product.** Research on pre-attentive processing confirms this: users form visual judgments about interfaces in 50ms (Lindgaard et al., 2006). Antaeus's four-layer model must feel instantaneous at every transition — Brief to Sheet should take under 100ms.

What Salesforce and HubSpot teach by negative example

The counter-examples are equally instructive. **Salesforce exposes its relational database structure directly to users** — objects, fields, related lists — rather than surfacing actions and insights. Users report that "Salesforce is a database and you are just filling in fields." (Quora) Its navigation features **10-20+ tabs**, Console versus Standard navigation modes, and a "More" dropdown that becomes a dumping ground. HubSpot spreads functionality across (Supportyourapp) **11 tabs and 6 Hubs**, each with its own sub-navigation. (Cargas |) For a founder with no sales team, **90% of HubSpot's features are irrelevant noise**. (Salesflare)

The specific anti-patterns to avoid:

- **Database-as-UI**: Surface actions and insights, not records and fields
- **Navigation proliferation**: Antaeus needs **≤5 primary navigation items**, not 16 modules
- **Feature-gating as business model**: All core features at one price, no premium-tier AI

- **Empty canvas without guidance**: Always answer "what's the single most important thing to do right now?"
- **Admin-dependent experience**: Self-service customization without gatekeeping
- **Complexity as lock-in**: Stickiness from value, not from switching costs

---

Fourteen psychological principles, ranked by scientific confidence and applicability

The behavioral psychology research reveals a hierarchy of principles by replication strength and relevance to Antaeus's specific use case. Here is the evidence base, synthesized for architecture decisions.

Tier 1: High confidence, high applicability

Implementation Intentions (Gollwitzer, 1999) have the strongest evidence base for Antaeus's use case. A meta-analysis of **94 studies and 8,000+ participants found a medium-to-large effect (d = 0.65)** on goal attainment. The mechanism: specifying "IF situation Y arises, THEN I will do Z" creates heightened cue accessibility and automatized response initiation. This directly validates Antaeus's "recommended next action" system. Every object surface should present a specific, contextual next action — not "follow up on this deal" but "Send the pricing proposal to Sarah Chen by Tuesday, since she mentioned budget approval happens Wednesdays." The if-then specificity is what drives the effect.

Goal Gradient Effect (Hull, 1932; Kivetz et al., 2006) is robustly replicated. Coffee shop customers purchased **more frequently as they approached a free coffee reward**, University of Chicago and the effect predicted customer retention. The "stuck in the middle" finding (Bonezzi et al., 2011) reveals a U-shaped motivation curve — high at start, dipping mid-goal, peaking near completion. **Antaeus's Readiness Score should be designed as a goal gradient accelerator**, with milestones at 25%, 50%, 75%, and 100% that each unlock a micro-celebration and reveal the next milestone.

Commitment and Consistency (Freedman & Fraser, 1966) shows that small initial commitments produce 4.5x compliance increases for larger follow-through requests. The onboarding implication: each micro-commitment (name your ICP → define 3 target accounts → draft one outreach message) creates escalating psychological investment. Heyflow The critical design constraint from research: the first ask must be **genuinely low-friction** — KlientBoost found that asking for company name as the first field decreased conversions by 50%. CXL

Cognitive Load Theory (Sweller, 1988) provides the theoretical foundation for the entire restructure. The distinction between extraneous load (design-imposed overhead, the enemy) and germane load (productive schema-building, the goal) Ballpark Medium explains why 16 flat

modules fail — every navigation decision between modules is extraneous load that contributes nothing to the user's sales effectiveness. **Hick's Law quantifies this: reducing visible navigation options from 16 to 4-5 should measurably decrease time-to-first-action.** LogRocket IxDF The Paradox of Choice research is nuanced — a 2010 meta-analysis found the average choice-overload effect "close to zero," To Summarise Varify but a 2015 meta-analysis of 99 studies identified that the effect IS significant when the domain is unfamiliar, options are complex, and consumers lack clear preferences. These conditions describe a technical founder's first encounter with a sales OS perfectly.

Tier 2: Strong evidence, requires careful application

Endowed Progress Effect (Nunes & Dreze, 2006) produced an **82% higher completion rate** (34% vs. 19%) when car wash loyalty cards started with 2 of 10 stamps pre-filled versus 0 of 8, despite identical effort required. Lead Alchemists For Antaeus onboarding: importing contacts, connecting email, and completing ICP setup should immediately register as "3 of 8 setup steps complete" — converting "not started" into "underway." The 10-25% initial endowment is the research-supported sweet spot. Coglode

Loss Aversion (Kahneman & Tversky, 1979) replicates at a **90% rate across 19 countries** Columbia University Mailm... (Ruggeri et al., 2020). Losses are psychologically ~2x as painful as equivalent gains. Behavioral Economics iResearchNet The IKEA Effect (Norton et al., 2012) compounds this — participants paid **63% more for items they assembled themselves.** Wikipedia ScienceDirect For Antaeus, this means making cross-object data flow visible is critical: when a user logs a call, show that this data appears on the contact timeline, the deal activity log, the account summary, and tomorrow's Brief. **The user must see their work has "tentacles" across the system.** The data portability paradox is also relevant: offering easy export paradoxically increases retention for strong products by functioning as a trust signal Policy Review (BEREC data shows 13.6% higher churn only in markets with mandated portability Clairfield — voluntary export signals confidence).

The Zeigarnik Effect requires the most nuanced treatment. A critical **2025 meta-analysis (Ghibellini & Meier, Nature Humanities & Social Sciences Communications) found NO general memory advantage for unfinished tasks**. However, it confirmed a robust **Ovsiankina Effect — a general tendency to resume interrupted tasks.** Nature nature The implication: don't rely on users *remembering* incomplete tasks (they won't); rely on the motivational drive to *resume* them, and surface incomplete states visibly in the UI. The research on loop closure versus loop transformation is decisive: cliffhanger studies (Schibler et al., 2023) found audiences desired future installments significantly more without sacrificing enjoyment. **The optimal pattern is "celebrate briefly → reveal next opportunity immediately."** Post-reward motivation drops sharply to baseline (Kivetz et al., 2006), making the moment after completion the highest churn risk. Coglode The Brief should never show "all done" — it should show "completed → here's what unlocked."

Tier 3: Applied carefully with ethical guardrails

Variable Ratio Reinforcement is the most powerful and most dangerous principle. Skinner's research shows variable ratio schedules produce the highest, most persistent response rates and are most resistant to extinction. (Medium) Eyal's Hook Model (Trigger → Action → Variable Reward → Investment) maps the mechanism to product design. (Medium) (Mindtools) The ethical distinction is critical: **variable rewards that create addiction (infinite scroll, loot boxes) exploit users, while variable insights that surface genuine patterns ("Your Tuesday emails get 3x higher open rates") serve users.** Antaeus should implement "the system notices your work and occasionally tells you something surprising" through pattern detection on user data, cross-object correlations, and anomaly alerts — not through intermittent notifications designed to maximize session time.

Self-Determination Theory (Deci & Ryan, 2000) identifies three needs: autonomy, competence, and relatedness. (Yu-kai Chou) The **overjustification effect** is the key risk — external rewards for intrinsically motivating activities undermine intrinsic motivation. (Simply Psychology +2) Points, badges, and leaderboards should be avoided entirely for core sales activities. Instead, provide competence-affirming feedback: "Your outreach response rate improved from 8% to 12% this week." For relatedness in a solo-user tool, research shows the tool itself can function as a relational entity through conversational AI, empathetic messaging, and "we understand founders like you" signals. **Anonymous stage-based benchmarking** ("Founders at your ARR typically have 23 active accounts") satisfies relatedness by connecting to an invisible cohort, but must include injunctive norms for outperformers to prevent the boomerang effect documented in the Opower study.

Peak-End Rule (Kahneman, 1993) has large confirmed effects across a 2022 meta-analysis of 174 samples. (PubMed Central) People judge experiences by their peak moment and ending, not their average or duration. (Wikipedia +2) **Every Antaeus session needs a designed peak (an insight delivery or achievement moment) and a designed end (progress summary + next session trigger).** Never end on an error, a timeout, or "nothing to show."

Temporal Discounting explains why stale deals don't feel urgent — the loss is in the future, and present tasks feel more pressing. (Medium) Visual aging indicators (green → yellow → orange → red) and explicit temporal framing ("Win rate drops 40% after 30 days of inactivity") make future costs salient now. But alarm fatigue is real: when everything is urgent, nothing is. **Signal decay should be selective, applying only to the 3-5 highest-impact items in the Brief.**

How every current module maps to the four-layer architecture

The restructure from 16 flat modules to a four-layer system requires each module to either dissolve into an object surface, become a system-level view, merge with another module, or be

eliminated. Here is the complete mapping, informed by every competitive pattern and psychological principle researched.

Layer 1 (Command): Brief / Grid / Queue

The Command layer absorbs **Dashboard, Signal Console, and Daily Cockpit** as its primary data sources, plus elements of **Readiness Score** (as the health/urgency ranking engine). The Brief generates a narrative operating summary (like Linear's project updates + Slack AI summaries). The Grid provides a rail + detail panel layout (like Linear's issue list + detail, or Attio's list view). The Queue provides a sequential triage flow (like Superhuman's Split Inbox).

What populates the Command layer: Ranked objects from across the system — signals requiring action, deals needing follow-up, accounts going stale, calls to prepare for, outreach sequences to review. Each item carries a state: **urgent** (red, decaying, loss-framed), **active** (blue, in-motion, momentum-framed), or **healthy** (green, compounding, no action needed). The ranking algorithm combines signal recency, deal value, days since last touch, stage-appropriate urgency thresholds, and AI-assessed risk.

Layer 2 (Sheet): Quick inspection

The Sheet rises from any Command mode item and absorbs elements of **every module that produces a record-level summary**. It shows: full object context (all attributes), gap visualization (missing fields framed as "win probability increases X% when complete"), connections to other objects (contact → account → deal → signals), the single recommended next action (implementation intention format), and downstream consequences ("if you do X, it flows to Y and Z").

Layer 3 (Workspace): Rail + morphing surface

This is where the most dramatic consolidation happens. The Workspace anchors on an object type and morphs its surface tabs based on that object's lifecycle needs:

Deal object surfaces (absorbing 6 modules):

- **Deal Workspace** → becomes the default Deal surface (pipeline position, value, contacts, timeline)
- **Discovery Studio** → becomes a "Discovery" tab on the Deal surface
- **Call Planner** → becomes a "Prep" tab on the Deal surface
- **Future Autopsy** → becomes a "Risk" tab on the Deal surface (pre-mortem analysis)
- **PoC Framework** → becomes a "Proof" tab on the Deal surface
- **Advisor Deploy** → becomes a "Team" tab on the Deal surface (advisor/champion mapping)

Account object surfaces (absorbing 3 modules):

- **Territory Architect v2** → becomes the Account surface's default view, with the hard 300 cap, solution-fit hypothesis, and swap mechanic
- **ICP Studio** → becomes a system-level configuration that manifests as "ICP Match" scoring on every Account surface
- **Sourcing Workbench** → becomes an "Add Accounts" action accessible from the Account list in Grid mode

Motion object surfaces (absorbing 3 modules):

- **Outbound Studio** → becomes the default Motion surface for email sequences
- **LinkedIn Playbook** → becomes a "LinkedIn" tab on the Motion surface
- **Cold Call Studio** → becomes a "Phone" tab on the Motion surface

System-level views (not object-specific):

- **Signal Console** → feeds into Layer 1 (Command) as the signal ingestion engine
- **Quota Workback** → becomes a system-level planning surface accessible via ⌘K or Settings
- **Playbook** → becomes the methodology backbone, surfacing contextual guidance on every object surface (like Salesforce's Path component, but non-blocking)
- **Readiness Score** → becomes the health scoring engine powering Layer 1 ranking and Layer 4 visualization
- **Methodology Hub** → merges with Playbook and becomes App Academy/FAQ/Help section

Layer 4 (Graph): Hidden diagnostic Absorbs **Readiness Score visualization** as a network graph showing object relationships and workspace health. Accessible via a discovery-moment button. This is where the IKEA Effect and Loss Aversion compound — the user sees the full web of their invested work, the "tentacles" of their data flowing across objects.

Modules that merge or transform

Current Module	Destination	Rationale
Dashboard	Eliminated — absorbed by Brief	Brief replaces dashboard with narrative intelligence
Tour	Eliminated — replaced by onboarding flow + contextual tooltips	Tours are extraneous cognitive load

Current Module	Destination	Rationale
Demo Lane	Transformed into guided first-run experience	Shows the product with seed data, not a separate mode
Welcome	Merged into onboarding flow	Single coherent first-run, not a separate screen
Settings	Remains as system-level, accessible via ⌘K	Standard utility surface

What each object type should make users feel

The research on Self-Determination Theory, pre-attentive processing, and peak-end rule converges on a clear design principle: **each object surface should evoke a specific emotional state, present one dominant action, and show what compounds automatically.**

ICP Object → Feel: Clarity and conviction. The founder should feel "I know exactly who I'm selling to." Dominant action: **Refine fit criteria.** Auto-compounds: ICP scoring flows to every Account, filtering the Brief and ranking the Queue. Surface shows the ICP definition with match distribution across the territory — a histogram of fit scores that makes the founder feel they've made a strategic decision, not just filled out a form.

Account Object → Feel: Strategic ownership. "This is MY territory — I chose these 300 and I know why." Dominant action: **Qualify or swap.** The 300-cap and swap mechanic create a zero-sum game that makes each account feel consequential. Auto-compounds: Account qualification status flows to Deal creation eligibility, Motion targeting, and Brief prioritization.

Signal Object → Feel: Time pressure with agency. "Something just happened, and I have a window to act." Dominant action: **Convert to motion.** Signals carry temporal decay (Clay's insight: common signals are noise, custom signals are alpha). Auto-compounds: Acting on a signal creates a Motion, which feeds the Queue, which builds pipeline. The decay visualization uses temporal discounting research — making the future cost of inaction feel present.
Sustainability-directory

Motion Object → Feel: Craftsmanship. "I'm building something personalized, not spamming." Dominant action: **Send next touchpoint.** Auto-compounds: Motion activity logs to Contact timeline, Deal activity, and Account last-touch — all visible as cross-object data flow. Outbound Studio, LinkedIn Playbook, and Cold Call Studio all dissolve into channel tabs on this surface.

Call Object → Feel: Preparation and confidence. "I know what I'm going to say and why." Dominant action: **Execute the call.** The Prep tab (formerly Call Planner) surfaces the account

context, deal stage, prior conversations, and recommended talk track. Auto-compounds: Call notes flow to Discovery surface on Deal, Contact timeline, and AI-generated follow-up tasks.

Deal Object → Feel: Momentum and control. "I can see exactly where this stands and what moves it forward." Dominant action: **Advance the stage.** Six former modules (Deal Workspace, Discovery Studio, Call Planner, Future Autopsy, PoC Framework, Advisor Deploy) become tabs that morph based on deal stage. Early stage emphasizes Discovery and Prep. Late stage emphasizes Proof and Team. Auto-compounds: Deal advancement updates pipeline metrics, Readiness Score, Brief narrative, and Quota Workback projections.

Proof Object → Feel: Evidence accumulation. "I'm building an irrefutable case." Dominant action: **Log evidence.** Auto-compounds: Proof items flow to Deal risk assessment (Future Autopsy surface) and advisor briefing materials.

Advisor Object → Feel: Leverage. "I'm not doing this alone — I have a strategic team." Dominant action: **Brief the advisor.** Auto-compounds: Advisor deployment logs to Deal timeline and informs the AI's recommended next action.

Handoff Object → Feel: Completion with legacy. "Everything I built transfers cleanly." Dominant action: **Package the handoff.** Auto-compounds: Handoff documentation flows from Deal history, Contact relationships, and Motion records — the full graph of the founder's work, visible and transferable.

Brief ranking logic: how urgency, signals, and decay models should work

The Brief is Antaeus's most critical surface — the first thing users see, the session's opening frame. Research on cognitive load, temporal discounting, and implementation intentions converges on a specific ranking architecture.

The ranking function should weight five factors:

1. **Temporal decay** (40% weight): Days since last touch relative to stage-appropriate thresholds. A deal in Discovery with no activity for 7 days is more urgent than an account with no activity for 7 days. Research on temporal discounting shows that explicit framing (Lead Alchemists) ("Win rate drops 40% after 14 days") makes abstract future loss feel present. Each object type should have calibrated decay curves — signals decay fastest (24-72 hours), deals decay by stage (Discovery: 5 days, Negotiation: 3 days), accounts decay slowest (14-30 days).
2. **Value magnitude** (25% weight): Deal value × win probability. Higher-value opportunities surface higher. This leverages pre-attentive processing — large bold numbers ("$47K at

risk") are processed before conscious attention, creating immediate emotional impact. (Daydreaming Numbers)

3. **Signal freshness** (15% weight): New signals (job changes, funding events, website visits) create time-limited windows. Clay's research shows common signals (funding) have shorter windows than uncommon signals (SOC2 certification). (Clay) Signal-specific decay rates prevent stale signals from cluttering the Brief.
4. **Gap severity** (10% weight): Missing critical data on high-value objects. A $100K deal with no champion identified should surface as "gap" urgency. Research on the Zeigarnik/Ovsiankina effect shows that visible incompleteness drives resumption behavior (Nature) (nature) — but only when framed as opportunity, not deficiency. "Identifying a champion increases win rate by 34%" rather than "No champion identified."
5. **Downstream impact** (10% weight): Objects whose action would cascade across the system. A call prep that would unlock a deal advancement, which would update pipeline metrics, which would affect Readiness Score — this cascading value should boost the item's rank.

The three density modes serve different cognitive states:

- **Brief** (narrative): For the first 30 seconds — "Here's what matters today" in 3-5 sentences. Minimal extraneous load. Peak-End Rule suggests this should open with the most positive or most impactful item. Generated by AI, similar to Linear's project updates combined with Slack's AI channel summaries.
- **Grid** (scannable): For the first 5 minutes — rail + detail panel showing all ranked items with enough context to decide. Like Linear's issue list with display options. Density toggle between compact (more items visible) and comfortable (more context per item), following Gmail's three-mode pattern.
- **Queue** (actionable): For execution mode — sequential processing of ranked items, Superhuman-style. Each item requires a disposition: Act (opens Sheet/Workspace), Snooze (set return time), Skip (next item), or Dismiss (remove from queue). Keyboard-driven: J/K to navigate, Enter to open, S to snooze, E to dismiss. (Blake Crosley)

Empty states that seed rather than stall

Empty states are the highest-risk moment in Antaeus's lifecycle. Research on the Endowed Progress Effect, Commitment and Consistency, and Cognitive Load Theory produces a specific strategy for each layer.

Command layer (Brief with no items): This should never happen after onboarding. The Brief should always have content. Before any user data: populate with methodology guidance ("Here's

what founders at your stage focus on first"), ICP configuration tasks framed as Brief items ("Define your first ICP segment — this becomes the filter for everything that follows"), and template deal objects that demonstrate the system. Attio auto-populates from email/calendar sync; Howtheygrow Clay provides dummy data in the Signals Hub. **Antaeus should combine both approaches: import real data where possible (email, calendar, LinkedIn) and fill gaps with guided template objects.**

Onboarding seeding of the Command layer: The onboarding flow should produce Brief items as a side effect of each step:

1. **Define ICP** → Brief item: "Review and refine your ICP match criteria" (active state)
2. **Import/create first 10 accounts** → Brief items: "Qualify each account for fit" (10 active items in Queue)
3. **Connect email** → Brief items: "Review contacts detected from email history" (active)
4. **Set first quota target** → Brief item: "Your Readiness Score is 15% — here's what moves it" (urgent)
5. **Complete first motion** → Brief item: "Follow up in 3 days" (scheduled, will become urgent)

By the end of onboarding (targeting **7 steps in under 5 minutes**, following Linear's radical simplicity model), the Brief should contain **8-15 ranked items**. candu The Endowed Progress Effect means showing "3 of 8 setup steps complete" immediately after account creation. SSRN The Goal Gradient Effect means designing milestones at 25%, 50%, 75%, and 100% that celebrate and immediately reveal the next phase. Ness Labs

Workspace layer (a deal surface with no discovery logged): This should show the gap as opportunity, not emptiness. Research on gap visualization across tools shows the effective pattern: "Discovery increases win rate by X%" with a single prominent CTA to start the discovery session. Pre-populate the Discovery tab with the recommended questions from Playbook methodology, showing them as a structured template Pitch (Pitch's template-first pattern) rather than a blank form. **The empty state IS the template** — it shows the user what "good" looks like and invites them to fill it in, not to build it from scratch.

Session design: the first 30 seconds, 5 minutes, and returning user

The Peak-End Rule, attention residue research, and temporal discounting converge on a specific session architecture.

First 30 seconds (new user, first session ever): The goal is competence (SDT) — the user must feel "I can do this" within 30 seconds. Following Linear's anti-onboarding philosophy: minimal steps, belief-system-first. candu Show the design thesis ("This is your strategic operating room")

for 3 seconds, then immediately drop into a guided ICP definition flow with **only one field to complete first** (following Commitment and Consistency: the first ask must be genuinely low-friction). After that single completion, show the Brief with the user's first item already ranked. **The peak moment of the first session should be the instant their first data appears in the Brief** — the "I built this" moment that activates the IKEA Effect. End the first session with a progress summary: "You've defined your ICP and started your territory. Tomorrow's Brief will show account qualification tasks."

First 5 minutes (any session): The entry is the Brief in narrative mode — 3-5 sentences covering what changed since last session, what's urgent today, and one variable-reward insight ("Deals where you scheduled demos within 48 hours close 2.7x faster — you have 2 opportunities to test this today"). Within 30 seconds, the user should be able to switch to Grid or Queue mode. **Cognitive load research demands that the first screen show ≤5 items** with clear visual hierarchy (pre-attentive processing: urgent items in saturated color, active items in standard weight, healthy items receded). The implementation intention pattern should appear on every item: not "Follow up with Acme" but "Send pricing proposal to Sarah at Acme — she mentioned budget review is Friday." Research shows this if-then specificity increases follow-through by d = 0.65.

Returning user (daily login): The Brief surface should function as a **"what changed since last session" summary** (like Linear's Inbox + Slack's catch-up). Three sections: (1) New signals requiring triage, (2) Deals that changed state or went stale, (3) Scheduled items that are now due. The variable reward is the insight — something the system noticed that the user couldn't have discovered alone. This creates the Hook Model's variable reward without the addictive pattern, because the reward is actionable intelligence, not engagement bait. Medium Amplitude **End every session with a designed closing moment**: "Today you advanced 2 deals, logged 3 calls, and your pipeline grew 8%. Tomorrow: prep for the Acme demo and follow up with 2 stale deals." This loads the next session's internal trigger (Eyal's Investment phase) and ensures the session ends positively (Peak-End Rule). Medium

---

How tension migrates forward and never resolves

The Zeigarnik model for Antaeus should implement **loop transformation, not loop closure**, based on the cliffhanger research (Schibler et al., 2023) showing that audiences desired future installments significantly more from cliffhangers without sacrificing enjoyment, and the post-reward reset phenomenon (Kivetz et al., 2006) showing motivation drops sharply after goal attainment. Coglode

The practical transformation chain for each object type:

Signal → Motion: Triaging a signal doesn't close the loop. It transforms into "Motion created — first touchpoint draft ready for review." The tension migrates from "something happened" to

"will they respond?" Celebrate the signal triage ("Signal captured → motion launched") for 1-2 seconds, then immediately surface the new open loop.

Motion → Contact response: Sending an outreach sequence doesn't close the loop. It transforms into "Sequence active — monitoring for engagement signals." The Brief shows "3 motions in flight, 1 showing engagement" — the open loops are visible and ranked by likelihood of response. Variable reward: the system alerts when someone opens the email or visits the pricing page (unexpected, useful, tied to the user's invested work).

Contact response → Deal creation: A positive response doesn't close the loop. It transforms into "Deal created — Discovery tab shows 0 of 6 recommended questions answered." The Endowed Progress Effect kicks in: the deal was created with the contact, account, and signal history already attached — "2 of 8 deal qualification steps complete" without the user doing deal-specific work. (SSRN) The Zeigarnik/Ovsiankina drive activates: the 6 unanswered discovery questions create motivational tension toward completion.

Deal advancement (any stage): Moving a deal from Discovery to Proposal doesn't close the loop. It transforms into "Proposal stage — PoC framework shows 0 of 4 proof points documented." The surface morphs: Discovery tab recedes, Proof tab becomes prominent. The Gap Gradient (combining Goal Gradient + gap visualization) shows "Deals at Proposal stage with documented proof points close 2.3x more often."

Deal won → Handoff: Winning a deal creates the highest-risk post-reward reset moment. Instead of showing "Congratulations, deal closed!" and stopping, the Brief immediately transforms: "Deal won — Handoff package 0% complete. Your first sales hire will need: account context, contact relationships, discovery notes, and competitive intelligence." The peak moment (deal won) is celebrated, but the end moment (session close) shows the next open loop.

The key ethical constraint: Every downstream loop must represent **genuine value** to the user, not artificial anxiety. "Document your discovery for the handoff" is genuinely useful — the founder will eventually hire and need this documentation. "Your Readiness Score dropped because you haven't logged enough calls this week" is punitive gamification. The distinction: loop transformation should reveal work that the user would want to do anyway, surfaced at the moment when motivation and context are highest.

Net-new surfaces the research suggests adding

Beyond the planned Daily Cockpit and App Academy, the competitive and psychological research suggests four additional surfaces:

1. "Since You've Been Gone" delta surface. Every tool studied that creates habitual return behavior has a change-detection layer. Linear's Inbox, Slack's catch-up, GitHub's notifications. Antaeus needs a dedicated delta computation that powers the Brief narrative. This isn't a

separate module — it's the engine beneath the Brief that compares current state to last-session state and produces ranked deltas.

2. Session close summary / "Loading the next trigger." Peak-End Rule research is unambiguous: the end of a session disproportionately affects remembered experience and return behavior. Antaeus should show a 3-5 second closing card when the user navigates away or closes the app: "Today: 3 calls logged, 2 deals advanced, pipeline +$28K. Tomorrow: Acme demo prep, follow up with 2 stale accounts." This is the Hook Model's Investment phase — loading the internal trigger for the next session.

3. Graph-powered "See What You Built" moment. Layer 4 exists in the plan, but the research suggests it should also be triggerable at milestone moments, not just on demand. When the user hits 50 accounts qualified, or $100K in pipeline, or 30 days of daily usage — the Graph surface should appear as a surprise reveal, showing the full web of their work. This is variable reward (unexpected, delightful) combined with IKEA Effect (seeing the full scope of labor invested) and Loss Aversion (making the switching cost visible and positive).

4. Stage-based behavioral benchmarks. No major sales tool currently provides in-product stage-based behavioral benchmarking for individual founders. This is a differentiation opportunity identified across the Social Proof research. "Founders at your ARR typically have 23 active accounts" provides the calibration that a first-time seller desperately needs. Carta's Total Compensation product proves the model at scale. The Opower study proves the mechanism works. **Include injunctive norms for outperformers** ("You're ahead of most founders at your stage ✓") to prevent the boomerang effect.

The architecture thesis, validated and refined

The research confirms the core Antaeus design thesis — "a strategic operating room, not a tool stack" — and provides the behavioral engineering to make it real. The four-layer model maps precisely to the four cognitive modes identified across all competitor analysis: **scan** (Brief), **evaluate** (Grid/Sheet), **execute** (Queue/Workspace), and **reflect** (Graph). Each layer has a specific psychological function:

- **Layer 1 (Command)** reduces extraneous cognitive load by pre-ranking decisions. It applies Hick's Law (fewer choices = faster decisions), temporal discounting (making future costs present), and implementation intentions (specific next actions).
- **Layer 2 (Sheet)** satisfies the competence need (SDT) by giving the user full context to make confident decisions. It applies progressive disclosure (essential context first, full detail on request) and gap visualization (non-punitive incompleteness).
- **Layer 3 (Workspace)** enables flow state (Csikszentmihalyi) by eliminating context-switching between modules. It applies attention residue research (Sophie Leroy: switching

tasks leaves cognitive residue that degrades performance by ~24%) and the object-centric navigation pattern (Attio, Linear: everything about an entity in one place).

- **Layer 4 (Graph)** delivers the peak moment (Peak-End Rule) and activates loss aversion + IKEA Effect by making the full web of invested work visible.

The transition from 16 flat modules to this four-layer system isn't just a UX reorganization — it's a **behavioral architecture** that engineers the session arc researched across every principle: low cognitive load at entry, escalating engagement through the middle, a designed peak, and a forward-loaded end that creates the internal trigger for return. Every object carries tension that migrates forward. Every surface shows the user what compounds. Every session ends with the next session's hook already loaded.

The product that results should feel like what Linear did to Jira, what Superhuman did to Gmail, and what Attio is doing to Salesforce — but applied to the specific, underserved problem of a technical founder who needs to sell before they can hire someone to sell for them.